

# Assignment4

Student number:1407870

2023-10-22

Subject name: A First Course In Statistical Learning

Subject code: MAST90104

---

# Q1

(a)

```
wine=read.csv('/Users/guanmuhan/Downloads/winequality-red.csv',sep = ';')
wine$quality = factor(wine$quality)
levels(wine$quality) <- c("bad","bad","average","average" , "good","good")
library(nnet)
mmod=multinom(quality~.,wine)
```

```
## # weights:  39 (24 variable)
## initial  value 1756.681050
## iter   10 value 1049.635007
## iter   20 value  689.921334
## iter   30 value  653.028931
## iter   40 value  651.960731
## iter   50 value  651.732058
## iter   60 value  651.695020
## iter   70 value  651.546418
## iter   80 value  651.305172
## iter   90 value  648.593999
## final   value  647.675831
## converged
```

```
summary(mmod)
```

```
## Call:
## multinom(formula = quality ~ ., data = wine)
##
## Coefficients:
##      (Intercept) fixed.acidity volatile.acidity citric.acid residual.sugar
## average   -178.8060   -0.26900223        -4.462724   -1.2742330   -0.259419065
## good      122.3626    0.05907234        -6.782318   -0.6097511    0.009571902
##      chlorides free.sulfur.dioxide total.sulfur.dioxide  density      pH
## average   -5.485951         0.01736726         0.0161405294  197.9551 -4.534165
## good     -13.977148         0.02721645        -0.0004499205 -119.7577 -3.941803
##      sulphates  alcohol
## average    1.199091  0.4310564
## good       5.012159  1.1267605
##
## Std. Errors:
##      (Intercept) fixed.acidity volatile.acidity citric.acid residual.sugar
## average    2.705097    0.1543634        0.835823    1.218232    0.08282364
## good       3.129535    0.1706778        1.110980    1.444455    0.09777426
##      chlorides free.sulfur.dioxide total.sulfur.dioxide  density      pH
## average    3.101619         0.02292630         0.008267819  2.646911  1.391239
## good       4.459550         0.02565747         0.009496124  3.052356  1.600228
##      sulphates  alcohol
## average    1.269845  0.1758899
## good       1.358037  0.1949868
##
## Residual Deviance: 1295.352
## AIC: 1343.352
```

```
mmodi=step(mmod,trace = F)
```

```
## trying - fixed.acidity
## trying - volatile.acidity
## trying - citric.acid
## trying - residual.sugar
## trying - chlorides
## trying - free.sulfur.dioxide
## trying - total.sulfur.dioxide
## trying - density
## trying - pH
## trying - sulphates
## trying - alcohol
## # weights: 36 (22 variable)
## initial value 1756.681050
## iter 10 value 962.548544
## iter 20 value 660.650129
## iter 30 value 652.962534
## iter 40 value 652.709099
## iter 50 value 652.619043
## iter 60 value 652.562298
## iter 70 value 652.480458
## iter 80 value 651.061574
## iter 90 value 648.331708
## final value 648.331656
## converged
## trying - fixed.acidity
## trying - volatile.acidity
## trying - citric.acid
## trying - residual.sugar
## trying - chlorides
## trying - total.sulfur.dioxide
## trying - density
## trying - pH
## trying - sulphates
## trying - alcohol
```

```
summary(mmodi)
```

```
## Call:
## multinom(formula = quality ~ fixed.acidity + volatile.acidity +
##          citric.acid + residual.sugar + chlorides + total.sulfur.dioxide +
##          density + pH + sulphates + alcohol, data = wine)
##
## Coefficients:
##          (Intercept) fixed.acidity volatile.acidity citric.acid residual.sugar
## average   -177.1713   -0.26252634          -4.565347   -1.4550010    -0.24141265
## good       125.7384    0.07207673          -6.952064   -0.8516209     0.02213188
##          chlorides total.sulfur.dioxide  density          pH sulphates  alcohol
## average   -5.434441          0.020732275  196.0641 -4.486351  1.255191  0.4420856
## good     -13.805282          0.006790962 -123.6146 -3.799881  5.050087  1.1345209
##
## Std. Errors:
##          (Intercept) fixed.acidity volatile.acidity citric.acid residual.sugar
```

```
## average      2.704340      0.1545768      0.8275218      1.201010      0.07916868
## good         3.125144      0.1708142      1.1010330      1.427641      0.09566638
##             chlorides total.sulfur.dioxide density      pH sulphates alcohol
## average      3.111854      0.006059025 2.649079 1.393743 1.273329 0.1756397
## good         4.468744      0.006894894 3.048687 1.597277 1.360258 0.1946937
##
## Residual Deviance: 1296.663
## AIC: 1340.663
```

*#We can see there are eleven beat\_is for each tow of the outcomes*

(b)

```
#First we convert quality into an ordered factor
wine$quality=ordered(wine$quality, levels=c("bad", "average", "good"))
library(MASS)
omod=polr(quality~.,data = wine)
summary(omod)
```

```
##
## Re-fitting to get Hessian
## Call:
## polr(formula = quality ~ ., data = wine)
##
## Coefficients:
##              Value Std. Error  t value
## fixed.acidity      0.098341   0.068831   1.42874
## volatile.acidity  -3.879626   0.544055  -7.13094
## citric.acid         0.053926   0.660585   0.08163
## residual.sugar      0.091284   0.054391   1.67830
## chlorides          -6.706752   1.966230  -3.41097
## free.sulfur.dioxide  0.004318   0.009655   0.44719
## total.sulfur.dioxide -0.003905   0.003244  -1.20383
## density            -86.533740   1.361587 -63.55357
## pH                 -1.523710   0.709843  -2.14654
## sulphates           2.683266   0.457355   5.86692
## alcohol             0.774485   0.078961   9.80847
##
## Intercepts:
##              Value Std. Error t value
## bad|average  -87.1941   1.3953  -62.4934
## average|good -80.5404   1.4091  -57.1567
##
## Residual Deviance: 1373.733
## AIC: 1399.733
```

```
omodi=step(omod,trace=F)
summary(omodi)
```

```
##
## Re-fitting to get Hessian
## Call:
## polr(formula = quality ~ volatile.acidity + chlorides + pH +
##       sulphates + alcohol, data = wine)
```

```
##
## Coefficients:
##               Value Std. Error t value
## volatile.acidity -4.0170    0.4688  -8.569
## chlorides        -6.5630    1.8940  -3.465
## pH               -2.1087    0.5221  -4.039
## sulphates         2.4779    0.4547   5.450
## alcohol           0.8712    0.0758  11.494
##
## Intercepts:
##               Value Std. Error t value
## bad|average  -3.0044   1.7780   -1.6898
## average|good  3.6146   1.7748    2.0366
##
## Residual Deviance: 1377.604
## AIC: 1391.604
```

There are much more parameters for the multinomial model than that for the ordinal model. Residual deviance higher for the ordinal model, and the df for the final ordinal model is 7 while the df for the final multinomial is 22. But the AIC for multinomial model is slightly smaller than that for the ordinal regression model, so we prefer to choose it.

(c)

```
#newdata for two models
newdata_multinomial=matrix(c(1,7.9,0.4,0.2,1.7,0.1,36,0.997,3.3,0.9,10),11,1)
newdata_ordinal=matrix(c(0.4,0.1,3.3,0.9,10),5,1)
#coefficient matrixes for two models
(coefficients_multi=matrix(coef(mmodi),2,11) )

##           [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,] -177.1713 -0.26252634 -4.565347 -1.4550010 -0.24141265 -5.434441
## [2,]  125.7384  0.07207673 -6.952064 -0.8516209  0.02213188 -13.805282
##           [,7]      [,8]      [,9]     [,10]     [,11]
## [1,]  0.020732275  196.0641 -4.486351  1.255191  0.4420856
## [2,]  0.006790962 -123.6146 -3.799881  5.050087  1.1345209

(coefficients_ordi=matrix(coef(omodi), 1,5) )

##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] -4.016974 -6.563039 -2.108722  2.477854  0.8712256

#linear combinations for two models
linear_predictors_multinomial = coefficients_multi%% newdata_multinomial
linear_predictors_ordinal=coefficients_ordi%% newdata_ordinal
#linear combinations for multinomial regression
linear_predictors_multinomial

##           [,1]
## [1,]  4.651586
## [2,]  2.365144

#linear combinations for ordinal regression
linear_predictors_ordinal

##           [,1]
## [1,]  1.720449
```

```

#inverse function for multinomial
softmax <- function(z) {
  exp_z <- exp(z)
  return(exp_z / sum(exp_z))
}
#probabilities for multinomial
probabilities_multinomial=softmax(c(0, linear_predictors_multinomial))
cat('By multinomial regression model, the probability of wine is bad:',probabilities_multinomial[1],'average:'

## By multinomial regression model, the probability of wine is bad: 0.008591316 average: 0.8999492 good: 0.101508684

#By ordinal regression
(g_r01=omodi$zeta[1]-linear_predictors_ordinal)

##           [,1]
## [1,] -4.724811

(g_r02=omodi$zeta[2]-linear_predictors_ordinal)

##           [,1]
## [1,] 1.894104

probabilities_ordinal_bad=1/(1+exp(-(g_r01)))
probabilities_ordinal_average=1/(1+exp(-(g_r02)))-probabilities_ordinal_bad
probabilities_ordinal_good=1-1/(1+exp(-(g_r02)))
cat('By ordinal regression model, the probability of wine is bad:',probabilities_ordinal_bad,'average:'

## By ordinal regression model, the probability of wine is bad: 0.00879436 average: 0.8604284 good: 0.130577256

```

(d)

$$odds\ ratio = \frac{\left[ \frac{r_{12}}{1-r_{12}} \right]}{\left[ \frac{r_{22}}{1-r_{22}} \right]} = \frac{e^{g(r_{12})}}{e^{g(r_{22})}} = e^{g(r_{12})-g(r_{22})} = e^{-(x_1-x_2)^T \beta} = e^{-(0.08-0.2)\beta_{chlorides}}$$

```

difference=-(0.08-0.2)
odds_ratio=exp(difference*omodi$coefficients[2])
cat('the odds ratio of ordinal regression model',odds_ratio)

## the odds ratio of ordinal regression model 0.4549514

```

## Q2

(a)

The likelihood is:

$$f(D|\beta) = \prod_i^n P(\epsilon_i|\beta) = (2\pi)^{-\frac{n}{2}} \exp\left(-\frac{1}{2} \sum_{i=1}^n (y_i - \beta x_i)^2\right)$$

(b)

We have Bayes theorem, the posterior is still normal

$$\text{posterior} \propto \text{likelihood} \times \text{prior}$$

$$P(\beta|D) \propto P(D|\beta)P(\beta) = \exp\left(-\frac{1}{2} \sum_i^n (y_i - \beta x_i)^2\right) \times \exp\left(-\frac{1}{200} \beta^2\right) = \exp\left(-\frac{1}{2} \left(\beta^2 \left(\sum_i^n x_i^2 + \frac{1}{100}\right) - 2\beta \sum_i^n y_i x_i + \sum_i^n y_i^2\right)\right)$$

A magnitude of  $\frac{1}{100}$  is markedly small and can often be disregarded in equations. To elaborate, this suggests that when the variance of the prior probability is substantial, the information it contributes to the estimation of posterior distribution is minimal. Under these circumstances, the estimation of the posterior distribution is main dominated by the likelihood.

Therefore, we can get mean and variance of posterior distribution of  $\beta$  by  $N(A, B)$ , where  $B = (\sum_i^n x_i^2)^{-1} = 0.04620728$ ,  $A = B(\sum_i^n y_i x_i) = 1.94572$

```
data=read.csv("/Users/guanmuhan/Downloads/simplereg (2).csv")
(B=1/sum(data$x^2))
```

```
## [1] 0.04620728
```

```
(A=B*sum(data$x*data$y))
```

```
## [1] 1.94572
```

(c)

```
model <- lm(y ~ 0+x , data=data) #no intercept
summary(model) #current beta is 1.94572
```

```
##
```

```
## Call:
```

```
## lm(formula = y ~ 0 + x, data = data)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max
```

```
## -1.6293 -0.4784 -0.1137  0.1198  1.5086
```

```
##
```

```
## Coefficients:
```

```
##      Estimate Std. Error t value Pr(>|t|)
```

```
## x      1.9457      0.1396   13.94 2.2e-14 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 0.6494 on 29 degrees of freedom
```

```
## Multiple R-squared:  0.8701, Adjusted R-squared:  0.8656
## F-statistic: 194.3 on 1 and 29 DF,  p-value: 2.195e-14
```

MH Algorithm

1. Initialize  $\text{beta\_current} = 1.94572$
2. FOR each iteration:
  - (a) Sample  $\text{beta\_new} \sim N(\text{beta\_current}, 0.25)$ .
  - (b) Compute  $\text{AR} = P(\text{beta\_new} \mid \text{data}) / P(\text{beta\_current} \mid \text{data})$ .  
(For normal distribution, the  $Q(\text{beta}/\text{beta}')/Q(\text{beta}'/\text{beta})$  is cancelled)
  - (c) Draw  $u \sim U(0,1)$ .
  - (d) IF  $u < \min(1, \text{AR})$ :  $\text{beta\_current} = \text{beta\_new}$ . Else:  $\text{beta\_current} = \text{beta\_current}$
  - (e) Record  $\text{beta\_current}$ .
3. take samples from the  $(T+1)$ -th iteration (throw out data generated in the burn-in stage)

(d)

```
# Provided parameters and data
beta_current <- 1.94572
sigma_prior <- sqrt(100) # std of priori dist
sigma_proposal <- 0.5    # std of proposal dist
iterations <- 20000     # number of iterations
burn_in <- 5000         # burn-in

#set seeds
set.seed(123)
# Metropolis-Hastings
beta_samples <- numeric(iterations - burn_in)
for (i in 1:iterations) {

  beta_new <- rnorm(1, mean = beta_current, sd = sigma_proposal)

  likelihood_current <- sum(dnorm(data$y - beta_current * data$x, mean = 0, sd = 1, log = TRUE))
  prior_current <- dnorm(beta_current, mean = 0, sd = sigma_prior, log = TRUE)

  likelihood_new <- sum(dnorm(data$y - beta_new * data$x, mean = 0, sd = 1, log = TRUE))
  prior_new <- dnorm(beta_new, mean = 0, sd = sigma_prior, log = TRUE)

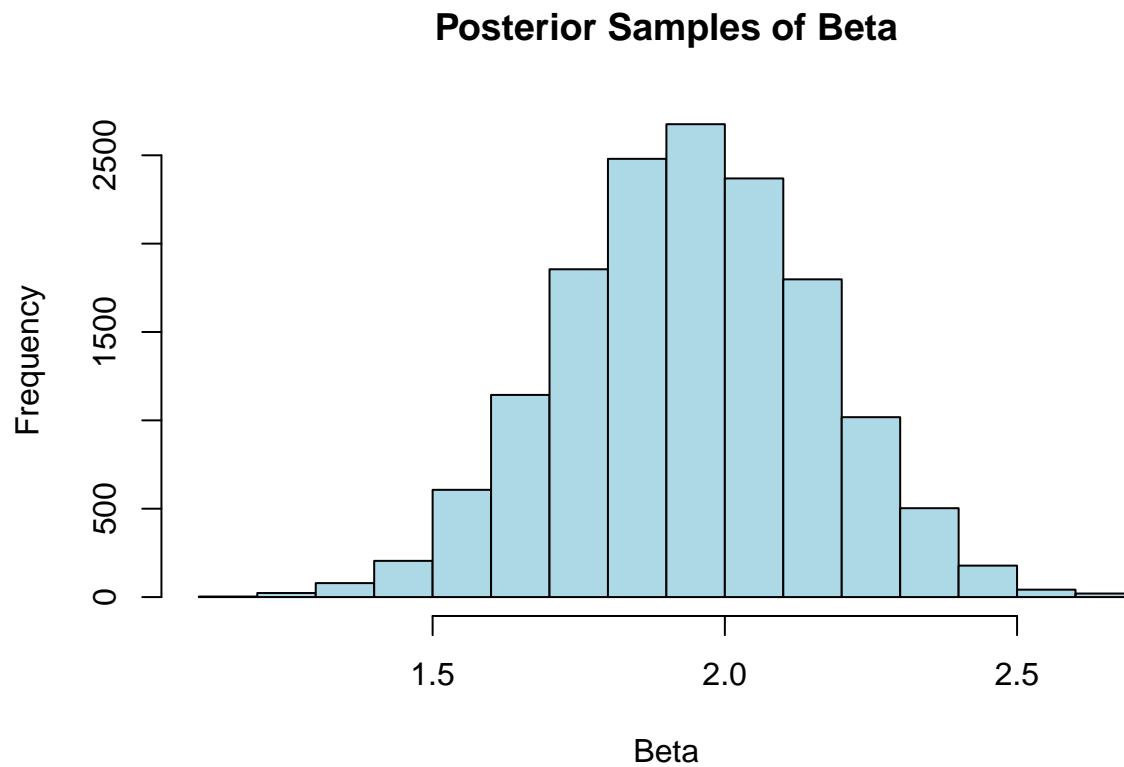
  acceptance_ratio <- exp((likelihood_new + prior_new) - (likelihood_current + prior_current))

  # Accept or reject new beta values
  if (runif(1) < acceptance_ratio) {
    beta_current <- beta_new
  }

  # Store sampling results (remove burn-in period)
  if (i > burn_in) {
    beta_samples[i - burn_in] <- beta_current
  }
}
```



```
hist(beta_samples, main = "Posterior Samples of Beta", xlab = "Beta", col = "lightblue", border = "black")
```



```
(sample_mean <- mean(beta_samples))
```

```
## [1] 1.940294
```

```
(sample_variance <- var(beta_samples))
```

```
## [1] 0.04789546
```

```
#By comparison ,the sampling distribution is close to the result in part b
```